

# Homework 1: Linear Regression

Due 23:59 on Monday, September 26, 2022

You will do this assignment individually and submit your answers as a PDF and code via the Blackboard course website. There is a mathematical component and a programming component to this homework.

## 1. MLE Estimate of the Bias Term (Bishop equation (3.19)) [10pts]

Let  $\Phi$  be our  $N \times J$  design matrix,  $\mathbf{t}$  our vector of  $N$  target values,  $\mathbf{w}$  our vector of  $J$  parameters, and  $w_0$  our bias parameter. As Bishop notes in (3.18), the sum-of-squares error function of  $\mathbf{w}$  and  $w_0$  can be written as follows

$$E(\mathbf{w}, w_0) = \frac{1}{2} \sum_{n=1}^N \left( t_n - w_0 - \sum_{j=1}^{J-1} w_j \cdot \phi_j(x_n) \right)^2.$$

Show that the value of  $w_0$  that minimizes  $E$  is

$$\begin{aligned} w_{0_{MLE}} &= \frac{1}{N} \sum_{n=1}^N t_n - \sum_{j=1}^{J-1} w_j \cdot \left( \frac{1}{N} \sum_{n=1}^N \phi_j(x_n) \right) \\ &= \bar{t} - \sum_{j=1}^{J-1} w_j \cdot \overline{\phi_j(x)} \quad \text{[compare to Bishop eqn. (3.19)]} \end{aligned}$$

## 2. Non-Uniformly Weighted Data [10pts]

Consider a data set in which each data point  $t_n$  is associated with a weighting factor  $r_n > 0$ , so that the sum-of-squares error function becomes

$$E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N r_n \{t_n - \mathbf{w}^\top \Phi(\mathbf{x}_n)\}^2.$$

Find an expression for the solution  $\mathbf{w}^*$  that minimizes this error function.

## 3. Priors and Regularization [10pts]

Consider the Bayesian linear regression model given in Bishop 3.3.1. The prior is

$$p(\mathbf{w} | \alpha) = \mathcal{N}(\mathbf{w} | \mathbf{0}, \alpha^{-1} \mathbf{I}),$$

where  $\alpha$  is the precision parameter that controls the variance of the Gaussian prior. The likelihood can be written as

$$p(\mathbf{t} | \mathbf{w}) = \prod_{n=1}^N \mathcal{N}(t_n | \mathbf{w}^\top \Phi(\mathbf{x}_n), \beta^{-1}),$$

Using the fact that the posterior is the product of the prior and the likelihood (up to a normalization constant), show that maximizing the log posterior (i.e.,  $\ln p(\mathbf{w} | \mathbf{t}) = \ln p(\mathbf{w} | \alpha) + \ln p(\mathbf{t} | \mathbf{w})$ ) is equivalent to minimizing the regularized error term given by  $E_D(\mathbf{w}) + \lambda E_W(\mathbf{w})$  with

$$E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (t_n - \mathbf{w}^\top \Phi(\mathbf{x}_n))^2$$
$$E_W(\mathbf{w}) = \frac{1}{2} \mathbf{w}^\top \mathbf{w}$$

Do this by writing  $\ln p(\mathbf{w} | \mathbf{t})$  as a function of  $E_D(\mathbf{w})$  and  $E_W(\mathbf{w})$ , dropping constant terms if necessary. Conclude that maximizing this posterior is equivalent to minimizing the regularized error term given by  $E_D(\mathbf{w}) + \lambda E_W(\mathbf{w})$ . (Hint: take  $\lambda = \alpha / \beta$ )

## 4. Modeling Motorcycle Helmet Forces [10pts]

The objective of this problem is to learn about linear regression with basis functions by modeling the g-forces associated with motorcycle helmet impacts. Download the file `motorcycle.csv` from the course website. It has two columns. The first one is the number of milliseconds since impact and the second is the g-force on the head. The data file looks like this:

```
"time since impact (ms)", "g force"
2.4, 0
2.6, -1.3
3.2, -2.7
3.6, 0
4, -2.7
6.2, -2.7
6.6, -2.7
6.8, -1.3
...
```

and you can see a plot of the data in Figure 1.

Implement basis function regression with ordinary least squares.<sup>1</sup> Some sample Python

---

<sup>1</sup>Note that the data clearly don't have fixed variance! There is obviously less variance on the left of the plot. Modeling such *heteroscedastic* data is beyond the scope of the course.




Figure 1: Motorcycle crash helmet data. The horizontal axis is time since impact and the vertical axis is force on the head.

code is provided in `linreg.py`. Plot the data and regression lines for the simple linear case, and for each of the following sets of basis functions:

- (a)  $\phi_j(x) = \exp\{ -((x - 10j)/5)^2 \}$  for  $j = 0, \dots, 6$
- (b)  $\phi_j(x) = \exp\{ -((x - 10j)/10)^2 \}$  for  $j = 0, \dots, 6$
- (c)  $\phi_j(x) = \exp\{ -((x - 10j)/25)^2 \}$  for  $j = 0, \dots, 6$
- (d)  $\phi_j(x) = x^j$  for  $j = 0, \dots, 10$
- (e)  $\phi_j(x) = \sin\{x/j\}$  for  $j = 1, \dots, 20$

In addition to the plots, provide one or two sentences for each, explaining whether you think it is fitting well, overfitting or underfitting. If it does not fit well, provide a sentence explaining why.

## 5. Explore Google Colab and Scikit-learn [10pts]

The objective of this problem is to gain some experience with Google Colaboratory and Scikit-learn.

1. Navigate to <https://colab.research.google.com/>. Complete the Welcome To Colaboratory tutorial. Nothing to submit for this. Just do the rest of the problem in Colab and submit the notebook.
2. In a new Colab notebook, load the diabetes regression dataset, see [https://scikit-learn.org/stable/datasets/toy\\_dataset.html](https://scikit-learn.org/stable/datasets/toy_dataset.html). Split the dataset into a training set, a validation set, and a test set.
3. Train linear regression [https://scikit-learn.org/stable/modules/linear\\_model.html](https://scikit-learn.org/stable/modules/linear_model.html) and  $k$ -nearest neighbors <https://scikit-learn.org/stable/modules/neighbors.html> models. Select the  $k$ -nearest neighbors  $k$  parameter using the validation set.
4. Report metrics that summarize algorithm performance on the test set [https://scikit-learn.org/stable/modules/model\\_evaluation.html](https://scikit-learn.org/stable/modules/model_evaluation.html).  $R^2$  and mean squared error are good metrics to include.

## 6. Calibration [0pt]

Approximately how long did this homework take you to complete?